

Input: **Q**: an atomic distributed priority queue that stores the requests to be scheduled. The priority is defined as follows: for **req i** and **req j**, the **P(i) > P(j)** if and only if **i < j** and **i** have loaded layers unexecuted. The subscription (e.g., **i**) indicates the request's arrival time.

At New Instance (target instance)

```
0  spawn(update_Q_when_layers_come) ▶ Run in the background
1  while execute under live:
2    |   q = Q.get_front()
3    |   forward_one_layer(q)
```

At Old Instance (source instance)

```
4  while execute under live:
5    |   ▶ Pull the pending requests with
6    |   If q, activation = pull_the_earlist_request(Q):
7    |   |   forward_all(q, activation)
```

