

1 req = receive() (a) vLLM
one score per instance
2 score = 4 * Q-BS + 1 * R-BS
3 sched_to = instances
4 .select_min(score)
5 req.forward(sched_to)

1 req = receive() (b) vLLM + KV\$-awareness
(e.g., Bailian-like)
2 kv_hit = KV\$.match(req)
3 score = $\lambda * (1 - kv_hit) + (1 - \lambda) * \text{norm}(BS)$
4 sched_to = instances
5 .select_min(score) Range [0,1]
6 req.forward(sched_to)

trim = 0.25cm YY.Ycm
XX.Xcm 0.25cm, clip