

**Per-iteration
Latency (ms)**




PyTorch




PyTorch w/ Mirage

1.9x 1.5x 1.0x




Chameleon-7B

1.4x 1.4x 1.4x




LLaMA-3-8B

1.2x 1.0x 0.9x



GPT-3-7B-LoRA

1.4x 1.4x 1.4x



nGPT-1B