

Self-Verification Hallucination

The agent reports verification success while judge spot checks expose unresolved artifact errors.

startUP-V1-015

51.52

Case score / 100
startUP-V1-015

1 Verification target

Membership anniversary workbook

Filter level-11 members and simulate route check-ins until each reaches level 12.

Inputs

member roster, route table, and store reward table

Rule

effective points = total points - expired points

Stop

stop immediately at the level-12 threshold

Output

XLSX with 8 required columns and exact names

Why this tests self-verification

A summary can look correct even when exact-name and cell-level spot checks fail.

2 Self-check drift

1

Local calculation succeeds

The agent computes thresholds and finds 30 level-11 members.

2

A plausible workbook is generated

The output has one table, 30 rows, and the requested 8 columns.

3

Summary is treated as validation

The final response restates the intended logic as if it were checked.

4

Rubric examples are not reopened

Exact names and current effective-point cells are not independently verified.

Model completion claim

"Completed: 30 members filtered; each route simulation stops at the level-12 threshold."

3 Claim vs judge evidence

Model self-check

Complete; 30 members; route simulation done.

Judge check

4 of 9 rubric items fail under spot checks.

Judge-inspected output rows

The model's own summary is positive, but artifact-level spot checks show mismatched cells.

Check	Case field	Model	Needed	Judge
Points	爱笑的猫猫大魔王	8497	7347	fail
Points	爱笑的龙猫君	8508	7548	fail
Points	爱笑的龙猫大叔	8527	7484	fail

Key miss: submitted H-column stores post-upgrade points, not current effective points.

Failure evidence

Names	Expected exact strings include spaces: 可爱的少年 1984, 小剑客 1984.
Points	Sampled effective-point cells are post-stop values, not current values.
Pattern	The final answer summarizes intended checks instead of rerunning judge-like checks.

Failed checks: 4 / 9

Score: 51.52 / 100

Why this supports the StartupBench story

Rubric-driven spot checks

The judge checks named members and exact cells rather than trusting a final summary.

Artifact-first grading

The benchmark reads the submitted XLSX table and compares values to expected outputs.

Self-check gap

The case separates generating a plausible file from verifying the file against acceptance examples.

Bad pattern takeaway

Self-verification hallucination occurs when the agent converts a progress summary into a pass signal. The missing step is independent artifact-level verification against rubric examples.