

Self-Verification Hallucination

The agent checks selected calculations, but does not reconcile the final workbook with source names and output-field semantics.

1 Verification target

Membership anniversary workbook

Filter level-11 members and simulate route check-ins until each reaches level 12.

Inputs member roster, route table, and store reward table

Rule effective points = total points - expired points

Stop stop immediately at the level-12 threshold

Output XLSX with 8 required columns and specified headers

Acceptance sensitivity

A summary can look correct even when exact-name and cell-level spot checks fail.

2 Self-check drift

1 Local calculation succeeds
The agent computes thresholds and finds 30 level-11 members.

2 A plausible workbook is generated
The output has one worksheet, 30 data rows, and 8 requested columns.

3 Verification coverage remains partial
The agent checks selected calculations and rows, but does not reconcile output fields with source semantics.

4 Artifact contradictions are missed
The reopened workbook stores post-upgrade points in the current effective-points column; this goes undetected.

Model completion summary (translated)

30 level-11 members were selected; each simulation stops when the level-12 threshold is reached.

3 Claim vs judge evidence

Model self-check

Complete; 30 members; route simulation done.

Judge check

Artifact checks expose unresolved name and field-value errors.

Judge-inspected output rows

The model's own summary is positive, but artifact-level spot checks show mismatched cells.

Check	Member code	Model	Needed	Judge
Points	Member M-01	8497	7347	fail
Points	Member M-02	8508	7548	fail
Points	Member M-03	8527	7484	fail

Submitted H-column stores post-upgrade points, not current effective points.

Failure evidence

Names Two source labels contain required spaces; submitted labels omit them.

Points Sampled effective-point cells are post-stop values, not current values.

Pattern Partial checks miss name and column-semantic mismatches.

StartupBench compares the submitted workbook with explicit task requirements.

Instruction-level spot checks

The judge checks named members and exact cells rather than trusting a final summary.

Artifact-first grading

The benchmark reads the submitted XLSX table and compares values to expected outputs.

Self-check gap

The case separates plausible file generation from instruction-level verification.

Pattern: overconfident self-verification. The agent reopens the workbook and checks selected calculations, but fails to reconcile exact names and the effective-points column with source data before declaring completion.